

Task inputs

 paper LaTeX source

☐ OpenReview
supplement

 verified artifact links

MATCH TARGET

metric · value ·
scope

no config given

ReAct tool loop any chat model, tools-only

WORKSPACE · LOGIN NODE, CPU

workspace_bash cwd-confined shell

write_file / edit_file author code & configs

update_plan running checklist

fetch_url public URLs only

CLUSTER · SLURM

list_partitions node pools via sinfo

run_gpu only path to a GPU:
salloc hold, srun steps

 billed by wall clock

wall clock metered vs the fixed H100-h grant

~500GB scratch · round cap + H100-h budget enforced by the harness

 **report.json** schema-constrained final pass

```
paper_id, claim, what_ran, scoring_command  
measurements[] {metric, observed_value,  
reference_value, scope, evidence}
```

```
agent_assessment ∈ {reproduced, partial,  
not_reproduced, could_not_run}
```

```
changes_made, blockers, evidence_files
```

 **evidence/** REPORT.md · commands.log ·
plan.md · run stdout

→ report + evidence

provenance auditor

grades report + evidence → **reproduced** / **partial** /
not_reproduced / **unverifiable**; echoed numbers
capped to 0