

Task inputs

 paper LaTeX source

☐ OpenReview supplement

 verified artifact links

MATCH TARGET

metric · value · scope

no config given

ReAct tool loop any chat model, tools-only

WORKSPACE · LOGIN NODE, CPU

workspace_bash cwd-confined shell

write_file / edit_file author code & configs

update_plan running checklist

fetch_url public URLs only

CLUSTER · SLURM

list_partitions node pools via sinfo

run_gpu only path to a GPU: salloc hold, srun steps, wall clock metered vs the H100-h grant

~500GB scratch · round + H100-h budgets harness-enforced

 **report.json** schema-constrained final pass

paper_id, claim, what_ran, scoring_command
measurements[] {metric, observed_value,
reference_value, scope, evidence}

agent_assessment ∈ {reproduced, partial,
not_reproduced, could_not_run}

changes_made, blockers, evidence_files

 **evidence/** REPORT.md · commands.log ·
plan.md · run stdout

→ report + evidence

provenance auditor

grades report + evidence → reproduced / partial /
not_reproduced / unverifiable; echoed numbers
capped to 0