

failure modes differ between weak and strong agents

	two weakest agents Qwen3.6-27B, MiniMax-M2.7	two strongest agents DeepSeek-V4-Flash, Muse Spark 1.2	change (pts)
Reimplemented but did not check the result	22.9%	10.9%	-12.0
Failed before any number was produced	13.3%	3.3%	-10.0
Build and dependency failures	10.1%	2.7%	-7.4
Wrong experiment	14.4%	10.3%	-4.0
Wrong artifact measured	9.0%	11.4%	+2.4
Never launched an experiment	3.2%	3.8%	+0.6
Echoed a shipped number	2.7%	2.2%	-0.5
All seven failure modes	75.5%	44.6%	-31.0