

NeurIPS 2025 submission · SharpZO

Match Target

Two-stage zeroth-order optimization improves prompt tuning for vision-language models using forward passes alone ...

metric test accuracy

scope EuroSAT test set

bar within tolerance

configuration withheld from the agent

Artifact Availability

Run data code weights execute or evaluate

Retrain data code weights train, then evaluate

Reimplement data code weights rebuild from the paper

Reproduction Agent

Reading the paper · Downloading artifacts · Writing code

Running experiments · Debugging · Reporting

workspace_bash write_file edit_file update_plan fetch_url list_partitions

run_gpu

run_gpu is the tool the agent calls to request GPU access, and it does the budget accounting for the run.

Provenance Audit

THE RUN CLAIMS

reproduced



THE AUDIT FINDS

partial

- ✓ each number traced to the execution that produced it
- ✗ protocol, split, and scale match the pinned configuration
- ✓ metric recomputed by the auditor from the raw outputs

flags scope-substitution medium



report + evidence

Run Output

```
[plan] locate released code, then run
$ git clone .../SharpZO
$ pip install -r requirements.txt
ERROR no match for flash-attn
[plan] retry with eager attention
$ run_gpu python eval.py --test
wrote evidence/eval.log
```

report.json

evidence/

REPORT.md plan.md
eval.log commands.log