

RECLAIM

can an agent reproduce a paper's central claim from what its authors released?

1

THE TASK



CQ-DINO: Mitigating Gradient Dilution via Category Queries...
NeurIPS 2025 · arXiv 2503.18430



Match target

metric

AP / AP50 / AP75

scope

V3Det val · 29,821 imgs

target

46.3 / 51.5 / 48.4 ±5%



configuration withheld



What the authors released

✓ data

✓ code

✓ weights



budget 8 H100-h

x100 papers

2

THE RUN



Reproduction agent

deepseek-v4

workspace_bash

write_file

fetch_url

 run_gpu

run 2889575 · 107 rounds · 6.0 of 8 H100-h

PLAN · r0-4

● [plan] AP 46.3 on V3Det · design eval, run it

FETCH · r5-8

○ \$ git clone github.com/.../CQ-DINO

○ \$ gdown ...swinb1k_v3det.pth # 1.89 GB

BUILD · r9-61

○ \$ pip install mmcv==2.2.0

● ERROR: no aarch64 wheel

○ \$ FORCE_CUDA=1 pip install -e mmcv/

RUN · r64

● \$ run_gpu tools/test.py # 29,821 imgs GPU

● srun: task 0: Out Of Memory

RESCUE & SCORE · r65-92

○ \$ cp /tmp/results.bbox.json evidence/

● \$ run_gpu chunked_eval.py # score them GPU

bbox_mAP 0.463 0.515 0.484 = paper's 46.3 / 51.5 / 48.4 ✓

● plan ○ step ● GPU run ● error \$ agent command

3

THE REPORT

report.json

claims: "reproduced"

evidence/ eval.log · commands.log

results.bbox.json · patches/

every number must trace to a logged run

4

THE AUDIT



Adversarial audit

re-derives the verdict from the evidence

✓ metric recomputed from raw predictions

✓ exact checkpoint, sibling rows ruled out

✓ AP_{S/M/L} values not in the paper: no echo

🚩 training arm untested (released ckpt only)

THE RUN CLAIMS reproduced

vs

THE AUDIT FINDS reproduced

RECLAIM scores the audited verdict;

across runs, claimed and audited diverge.