

RECLAIM

one peer-reviewed paper per task

1 TASK

match_target

metric test accuracy
scope EuroSAT test set
bar within tolerance
config withheld

Artifact availability

	code	data	wtg
Run	✓	✓	✓
Retrain	✓	✓	✗
Reimplement	✗	✓	✗

the tier records how much of the result the agent must rebuild

availability cliff

reproduction falls away as the release thins

2 RUN

workspace_bash write_file

edit_file update_plan

fetch_url list_partitions

run_gpu

```
[plan] locate released code
$ git clone .../SharpZ0
$ pip install -r reqs.txt
ERROR no match: flash-attn
[plan] retry, eager attn
$ run_gpu python eval.py
wrote evidence/eval.log
```

report.json

evidence/

REPORT.md plan.md
eval.log commands.log

A ReAct loop over a workspace shell and a metered Slurm session. run_gpu is the only path to a GPU and its wall clock is charged against the run's grant.

3 AUDIT

THE RUN CLAIMS

agent_assessment:
reproduced

- ✓ each number traced to its execution
- ✗ protocol, split, and scale match the pinned config
- ✓ metric recomputed from raw outputs

flags scope-substitution
medium

THE EVIDENCE SUPPORTS
partial

high-severity flag → disqualified
artifact wall → capped at blocked

self-claim gap

what the run claims against what its evidence supports